



Machine Learning and Optimisation

COMP24111 Course Notes

Jonathan Tang

Copyright © 2017 Jonathan Tang

[jtang0506.github.io](https://github.com/jtang0506)

Contents

I	Section I	
1	Introduction to Machine Learning	7
1.1	Supervised learning	7
1.1.1	Classification	7
1.1.2	Regression	7
1.2	Unsupervised learning	7
1.3	Reinforcement learning	8
2	k-nearest Neighbour	9
2.1	k-nearest Neighbour Classification	9
2.2	k-Nearest Neighbour Regression	10
3	Linear Classification and Regression	11
4	Logistic Regression	13
5	Support Vector Machine	15
5.1	Model Validation	15
5.1.1	Holdout method	15
5.1.2	Random subsampling	15
5.1.3	k-fold cross validation	15
5.1.4	Leave-one-out cross validation	15
5.1.5	Bootstrap	16
5.2	Confusion matrix	16

6	Deep Learning Models	19
----------	-----------------------------	-----------

II

Section I

7	Generative Models and Naive Bayes	23
8	Clustering Analysis Basics	25
8.1	Introduction	25
8.2	Data Representation	25
8.3	Distance Measures	26
9	k-means Clustering	29
9.1	Partitioning clustering approach	29
9.2	Introduction	29
9.3	k-means algorithm	29
9.4	Issues	30
10	Hierarchical and Ensemble Clustering	31
10.1	Hierarchical clustering approach	31
11	Cluster Validation	33
	Index	35



Section I

1	Introduction to Machine Learning	7
1.1	Supervised learning	
1.2	Unsupervised learning	
1.3	Reinforcement learning	
2	k-nearest Neighbour	9
2.1	k-nearest Neighbour Classification	
2.2	k-Nearest Neighbour Regression	
3	Linear Classification and Regression	11
4	Logistic Regression	13
5	Support Vector Machine	15
5.1	Model Validation	
5.2	Confusion matrix	
6	Deep Learning Models	19



1. Introduction to Machine Learning

1.1 Supervised learning

In **supervised learning** problems, there is an input, X , an output, Y , and the task is to learn the relationship between the input and the output. A training example in supervised learning is the pair (x, y) where x is the input and y is the target output. We assume a model defined up to a set of parameters, $y = g(x | \theta)$ where $g(\cdot)$ is the model and θ are its parameters.

1.1.1 Classification

For a classification problem, the task of the classifier is to assign a class label to a given input.


1.1.2 Regression

Suppose we want to predict the salary of a Computer Science graduate role. Inputs are the students attributes - degree result, previous experience and other information - that we believe affects a graduate's worth. The output is the predicted salary. Such problems where the output is a continuous number is known as a **regression** problem.

For our model, $y = g(x | \theta)$, y is a number since this is a regression model. In the simple example above, the model optimises the parameters, θ , such that the approximation error is minimised. For out example, the model is linear with the form $y = wx + w_0$ where w and w_0 are the parameters optimised for best fit to the training data.

1.2 Unsupervised learning

Unlike in supervised learning, we do not have a supervisor and we only have input data. The task in **unsupervised learning** is to form a natural understanding of the hidden structure of unlabelled data. There is a structure to the input space such that certain patterns occur more often than others, and we want to see what generally happens and what does not.



1.3 Reinforcement learning

2. k-nearest Neighbour

The **k-nearest Neighbour** estimation is one of the simplest machine learning algorithms, it is a **non-parametric method** (no parameter needs to be optimised) and requires **no explicit training**.

2.1 k-nearest Neighbour Classification

The *k*-nearest neighbour classifier assigns an instance to the class most heavily represented among its *k* neighbours. It is based on the idea that the more similar the instances, the more likely it is that they belong to the same class. We measure how similar two data points are by using a reasonable similarity or distance measure.

Definition 2.1.1 — Euclidean Distance. Given two *d*-dimensional data points, **p** and **q**, where $\mathbf{p} = [p_1, p_2, \dots, p_d]$ and $\mathbf{q} = [q_1, q_2, \dots, q_d]$, we define the **Euclidean distance** $d(\mathbf{p}, \mathbf{q})$ as:

$$d(\mathbf{p}, \mathbf{q}) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \dots + (p_d - q_d)^2} = \sqrt{\sum_{i=1}^d (p_i - q_i)^2}$$

In this case, the *k*-nearest neighbours are the training points with the lowest Euclidean distances to the testing point.

Definition 2.1.2 — Inner Product. Given two *d*-dimensional data points, **p** and **q**, where $\mathbf{p} = [p_1, p_2, \dots, p_d]$ and $\mathbf{q} = [q_1, q_2, \dots, q_d]$, we define the **inner product** as:

$$s_{inner}(\mathbf{p}, \mathbf{q}) = \sum_{i=1}^d p_i q_i$$

Since the inner product is a similarity measure, the *k*-nearest neighbours are the training points possessing the highest similarity values to the testing point.

2.2 k-Nearest Neighbour Regression



3. Linear Classification and Regression



4. Logistic Regression



5. Support Vector Machine

5.1 Model Validation

5.1.1 Holdout method

With the **holdout method**, the whole dataset is split into two groups, a training set and a testing set. The model is trained using the training set and asked to predict output values for the data in the testing set, which it has never seen before. This evaluation can have a high variance as we may get an unfortunate split as it may depend heavily on which data points end up in the training set and which end up in the testing set. Another drawback is that we may only have a small dataset, so we may not afford to set aside a portion of the dataset for testing.

5.1.2 Random subsampling

5.1.3 k-fold cross validation

k-fold cross validation is one way to improve over the holdout method, the dataset is divided into k partitions and the holdout method is performed k times. Each time, one of the k subsets is used as the testing set and the remaining $k - 1$ subsets are combined and used as the training set. We evaluate the error estimate as the average error across the k trials. The advantage of this method is that it matters less how the dataset gets divided and all the examples in the dataset are eventually used for both training and testing, therefore the variance of the resulting estimate is reduced as k is increased. However, the training has to be done k times, which means this method gets expensive to compute as k is increased.

5.1.4 Leave-one-out cross validation

Leave-one-out cross validation is the degenerate case for k -fold cross validation, with $k = N$, the number of data points in the dataset. As before, the average error is computed and used to evaluate the model. The evaluation given by leave-one-out cross validation error is good, but it is very expensive to compute compared to other methods.

5.1.5 Bootstrap

5.2 Confusion matrix

A confusion matrix is a table with two rows and two columns which allows further analysis rather than just the proportion of correct classifications. For example, if we had a dataset containing 98 examples from Class 1 and 2 examples from Class 2, there may be a classifier which classifies all the observations as Class 1, which yields a overall accuracy of 98%. However, accuracy is not a reliable metric for this classifier as the classifier would have 100% recognition rate for Class 1 but 0% recognition rate for Class 2.

■ **Definition 5.2.1 — True Positives (TP).** Both the predicted class and actual class is 'Yes'.

■ **Definition 5.2.2 — True Negatives (TN).** Both the predicted class and actual class is 'No'.

■ **Definition 5.2.3 — False Positives (FP).** The predicted class is 'Yes' and actual class is 'No'.

■ **Definition 5.2.4 — False Negatives (FN).** The predicted class is 'No' and actual class is 'Yes'.

Worked Example

Suppose that we predicted the presence of a disease that a patient and obtain the results below.

n = 165		Predicted Class	
		No	Yes
Actual Class	No	50	10
	Yes	5	100

■ **Example 5.1** For the above example, **TP** = 100, **TN** = 50, **FP** = 10 and **FN** = 5. ■

Some questions we may want to ask about our predictions are:

1. When the patient actually has the disease, how often did we predict this?
2. When the patient does not have the disease, how often did we predict this?
3. When we predict that a patient has a disease, how often are we correct?

First, we define some terms which directly relate to the questions above.

■ **Definition 5.2.5 — Sensitivity / Recall.** The proportion of positives that are correctly classified.

$$\frac{\text{True Positives}}{\text{Number of Actual Positives}} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

■ **Definition 5.2.6 — Specificity.** The proportion of negatives that are correctly classified.

$$\frac{\text{True Negatives}}{\text{Number of Actual Negatives}} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

■ **Definition 5.2.7 — Precision.** The proportion of predicted positives which were positive.

$$\frac{\text{True Positives}}{\text{Number of Predicted Positives}} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

■ **Example 5.2** By definition, the answers to the above questions are sensitivity, specificity and precision respectively. ■

Exercise 5.1 Calculate the sensitivity, specificity and precision for the confusion matrix given above. ■



6. Deep Learning Models



Section I

7	Generative Models and Naive Bayes	23
8	Clustering Analysis Basics	25
8.1	Introduction	
8.2	Data Representation	
8.3	Distance Measures	
9	k-means Clustering	29
9.1	Partitioning clustering approach	
9.2	Introduction	
9.3	k-means algorithm	
9.4	Issues	
10	Hierarchical and Ensemble Clustering	31
10.1	Hierarchical clustering approach	
11	Cluster Validation	33
	Index	35



7. Generative Models and Naive Bayes

8. Clustering Analysis Basics

8.1 Introduction

Clustering analysis is the process of finding similarities between data according to the characteristics underlying the data and grouping similar data objects into clusters. A **cluster** is a group of data points which are similar to one another within the same group and dissimilar to the points in other groups. Clustering analysis is **unsupervised** as there are no predefined classes for a training data set. Typically, clustering analysis is either used as a stand-alone tool to gain an insight into data distribution or used as a preprocessing step of other algorithms in intelligent systems.

The two general tasks of clustering analysis are:

1. Identifying the "natural" number of clusters present in a data set
2. Properly grouping data points into "sensible" clusters

Real Life Applications

- **Marketing** - companies can discover distinct groups in their customer bases and then use this knowledge to develop targeted marketing programs
- **Social network mining** - discovering communities of similar interests in a large group of people
- **Image segmentation** - dividing an image into distinct regions for object recognition

8.2 Data Representation

A data matrix is a $n \times p$ matrix which represents n data points with p dimensions. The data matrix has *two* modes as rows and columns represent different entities.

$$\begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1p} \\ x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \vdots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

A distance / dissimilarity matrix is a $n \times n$ matrix which represents the distance between each data point. The distance matrix is a symmetric triangular matrix as $\mathbf{d}(x, y) = \mathbf{d}(y, x)$ where x, y are two data points. The distance matrix has *one* mode as the rows and column represent the distance for the same entity.

$$\begin{bmatrix} 0 & & & & \\ d(x_2, x_1) & 0 & & & \\ d(x_3, x_1) & d(x_3, x_2) & 0 & & \\ \vdots & \vdots & \vdots & \ddots & \\ d(x_n, x_1) & d(x_n, x_2) & d(x_n, x_3) & \cdots & 0 \end{bmatrix}$$

8.3 Distance Measures

The **Minkowski distance** is a metric in a normed vector space which can be considered as the generalisation of both the Euclidean distance and the Manhattan distance.

Definition 8.3.1 — Minkowski Distance. The Minkowski distance of order p between two points $\mathbf{x} = (x_1, x_2, \dots, x_n)$ and $\mathbf{y} = (y_1, y_2, \dots, y_n)$ is defined as:

$$d(\mathbf{p}, \mathbf{q}) = \left(|x_1 - y_1|^p + |x_2 - y_2|^p + \cdots + |x_n - y_n|^p \right)^{\frac{1}{p}}, p > 0$$

The value of p in the Minkowski Distance would be selected depending on the application. The **Manhattan** and **Euclidean** distances are simply Minkowski distances of order 1 and 2 respectively.

Definition 8.3.2 — Manhattan Distance.

$$d(\mathbf{p}, \mathbf{q}) = |x_1 - y_1| + |x_2 - y_2| + \cdots + |x_n - y_n|$$

This is also known as the city block distance, the reason becomes clear in the examples below.

Definition 8.3.3 — Euclidean Distance.

$$d(\mathbf{p}, \mathbf{q}) = \sqrt{|x_1 - y_1|^2 + |x_2 - y_2|^2 + \cdots + |x_n - y_n|^2}$$

This is the distance that you are already familiar with, in n dimensions.

Exercise 8.1 When is the Manhattan distance between x and y equal to the Euclidean distance between x and y , where x and y are data points of two dimensions? ■

Definition 8.3.4 — Cosine measure. For $\mathbf{x} = (x_1, x_2, \dots, x_n)$ and $\mathbf{y} = (y_1, y_2, \dots, y_n)$

$$\cos(\mathbf{x}, \mathbf{y}) = \frac{x_1 y_1 + \cdots + x_n y_n}{\sqrt{x_1^2 + \cdots + x_n^2} \sqrt{y_1^2 + \cdots + y_n^2}} \quad \text{and} \quad d(\mathbf{x}, \mathbf{y}) = 1 - \cos(\mathbf{x}, \mathbf{y})$$

Note that $-1 \leq \cos(\mathbf{x}, \mathbf{y}) \leq 1$ and $0 \leq d(\mathbf{x}, \mathbf{y}) \leq 2$

Worked Example

Suppose we have some documents that we want to compare to see how similar texts are, without taking into account the order of the words. For this example, we will consider:

- (a) Labeeba loves me more than Hani loves me
- (b) Hani loves Labeeba more than Labeeba loves Hani
- (c) Labeeba likes Hani more than Hani likes Labeeba

First, we make a list of all the words that appear across all texts.

[Hani, Labeeba, likes, loves, me, more, than]

Next, we construct a vector for each document, with the frequency of each word.

$$\mathbf{a} = [1, 1, 0, 2, 2, 1, 1]$$

$$\mathbf{b} = [2, 2, 0, 2, 0, 1, 1]$$

$$\mathbf{c} = [2, 2, 2, 0, 0, 1, 1]$$

Now we can compute the **cosine measures**: $\cos(\mathbf{a}, \mathbf{b})$, $\cos(\mathbf{a}, \mathbf{c})$ and $\cos(\mathbf{b}, \mathbf{c})$.

$$\cos(\mathbf{a}, \mathbf{b}) = \frac{(1)(1) + (1)(2) + (0)(0) + (2)(2) + (2)(0) + (1)(1) + (1)(1)}{\sqrt{1^2 + 1^2 + 0^2 + 2^2 + 2^2 + 1^2 + 1^2} \sqrt{2^2 + 2^2 + 0^2 + 2^2 + 0^2 + 1^2 + 1^2}} \approx 0.694$$

By similar calculations, we find that $\cos(\mathbf{a}, \mathbf{c}) \approx 0.463$ and $\cos(\mathbf{b}, \mathbf{c}) \approx 0.714$. This tells us that sentences **b** and **c** are the most similar pair out the above texts, as they have the highest cosine measure. This is also the reason $d(\mathbf{x}, \mathbf{y})$ is defined as $1 - \cos(\mathbf{x}, \mathbf{y})$ since a larger cosine measure would be converted into a smaller distance.



9. k-means Clustering

9.1 Partitioning clustering approach

A partitioning clustering approach is done by iteratively partitioning the training data set to learn a partition of the given data space to produce several non-empty clusters, where the number of clusters is usually given in advance. In principle, an optimal partition is achieved by minimising the sum of the squared distance to its "representative object" in each cluster.

Exercise 9.1 Why do we have to use the squared distance? ■

9.2 Introduction

The k-means algorithm is the simplest partitioning method for clustering analysis and is widely used in data mining algorithms. It is a **heuristic method** and is based on each clustering being represented by the centre of the cluster and the algorithm will converge to stable centroids of clusters. The goal of this algorithm is to find a partition of k clusters to optimise the chosen partition criterion - the global optimum is achieved by exhaustively searching all partitions.

9.3 k-means algorithm

Before we can use the k-means algorithm, we must decide the cluster number k and an appropriate distance measure that will be used. We initialise the algorithm by selecting k distinct seed points, randomly.

1. assign each object to the cluster of the nearest point
2. compute new seed points as the centroids of the clusters of the current partition

The two steps are repeated until it converges (membership in each cluster no longer changes). At this point, we have k centroids which partition the whole data space into k mutually exclusive subspaces to form a partition.




Figure 9.1: Voronoi Diagram

The figure above shows how a data space can be partitioned, where each black dot represents a centroid of a cluster. Notice that the distance used above is Euclidean distance.

9.4 Issues

- This algorithm has computational complexity of $O(tKn)$, where n is the number of data points, K is the number of clusters and t is the number of iterations. Normally, $K, t \ll n$.
- It is sensitive to initialisation, which may result to unwanted solutions
- Is unable to handle noisy data and outliers, which results in an inaccurate partition
- Requires prior knowledge of the cluster number
- Incapable of handling clusters of non-convex shape
- Inapplicable to categorical data, since the mean is not defined
- How do we evaluate the k-mean performance?

Example

The k-Means algorithm is very simple, so an example has been omitted. An example may be added in the future, if I have spare time - but for now, just check the lecture slides.




10. Hierarchical and Ensemble Clustering

10.1 Hierarchical clustering approach

Hierarchical clustering approach is done by sequentially partitioning the data set to construct nested partitions layer by layer, via grouping objects into a tree of clusters. We use a generalised distance matrix as the clustering criteria and we do not need to know the number of clusters in advance.

Approaches to hierarchical clustering

There are two main approaches to hierarchical clustering: **agglomerative clustering** (bottom-up) and **divisive clustering** (top-down).



11. Cluster Validation



Index

- agglomerative clustering, 31
- centroid, 29
- cluster, 25
- confusion matrix, 16
 - false negative, 16
 - false positive, 16
 - precision, 16
 - recall, 16
 - sensitivity, 16
 - specificity, 16
 - true negative, 16
 - true positive, 16
- data representation, 25
- distance measures, 26
 - Cosine measure, 26
 - Euclidean distance, 26
 - Manhattan distance, 26
 - Minkowski distance, 26
- divisive clustering, 31
- hierarchical clustering, 31
 - approaches, 31
- k-nearest Neighbour, 9
 - classification, 9
 - regression, 10
- model validation, 15
 - holdout method, 15
 - k-fold cross validation, 15
 - LOO cross validation, 15
- partitioning clustering, 29
- reinforcement learning, 8
- similarity / distance measures
 - Euclidean distance, 9
 - inner product, 9
- Supervised learning
 - regression, 7
- supervised learning, 7
 - classification, 7
- unsupervised learning, 7